

HoneyPot Fleet — Monthly Threat Report, May 2026

TLP:CLEAR · Reporting period: 2026-05-01 — 2026-06-01 (UTC)

Executive summary

Over May 2026 our distributed honeypot sensors (deployed in Germany, the United States and Russia) recorded roughly 911,000 attack events across SSH, multi-protocol exploitation, Docker, LLM, MikroTik and Kubernetes surfaces. The headline development of the month is the continued maturation of **attacks against self-hosted LLM inference endpoints** (Ollama, llama.cpp): a distinct attack category with its own taxonomy, tooling and persistent operators — now extending beyond compute theft into **SSRF-based theft of cloud credentials**. Other notable activity includes ongoing targeting of Solana validators via SSH, multi-architecture Mirai/Gafgyt droppers, a dominant WannaCry/SMB stream in captured malware, and the first ICS interactions (IEC-104 OLTC commands) against a newly deployed substation decoy.

Volume summary

SSH (Cowrie), exact counters:

Metric	DE	US	RU	Total
Sessions	155,620	48,240	228,858	433,871
Successful logins	3,876	525	4,280	8,753
Failed logins	121,842	11,837	156,415	290,621
Shell commands	1,413	184	1,483	3,080
Payload downloads	84	1	0	85

Fleet-aggregate service totals: Dionaea (SMB/MSSQL/MySQL/FTP/SIP) ~734,000 events; Docker API 4,296; Ollama (LLM) ~75,300 HTTP requests; Winbox (MikroTik) 5,700; Kubelet 2,006. Russia leads SSH volume but shows the lowest success rate (2.7%) — much of the RU traffic is blind brute-forcing whose dictionaries fit the decoy credentials poorly — and received zero payload downloads: attackers perform reconnaissance (uname) but do not proceed to delivery.

Key finding 1 — Attacks against LLM inference endpoints (headline)

The fake Ollama API received **~75,300 requests** in May. Distribution of the classified portion by activity: ~95% scanning, ~3.7% inference abuse (free GPU/compute theft via `/api/generate`, `/v1/chat/completions`), ~0.8% external model abuse, and — new this month — **~0.5% `ssrf_cloud_cred`**: use of the LLM endpoint as an SSRF pivot toward cloud metadata services (169.254.169.254) to steal IAM credentials. This marks a qualitative shift from compute theft to compromise of the underlying infrastructure.

Targeting is not limited to Ollama: the `/props` endpoint (1,282 requests) is specific to **llama.cpp** servers, and adjacent **n8n** probing (`/rest/oauth2-credential/callback`, `/api/v1/webhook/test`) indicates hunting for AI-orchestration credentials. The dominant scanning tool is the **official Ollama Go client** (UA `ollama/0.0.0 (amd64 linux) Go/go1.26.3`, 50,529 requests), alongside the commercial AI-asset scanner EchelonGraph-ShadowAIRadar.

Longitudinal data shows a **persistent operator**: 203[.]159[.]90[.]152 (AS210558 1337 Services GmbH, a network associated with the Stark Industries ecosystem per open sources) dominated this surface in March–April and remained active in May (~25% of all requests to the decoy). The May leader 37[.]19[.]210[.]21 (AS212238 Datacamp/CDN77, ~67% of requests) is likely a VPN exit fronting the real operator. Note: AS212238 is a large shared hosting provider and appears for multiple unrelated tenants in this report.

Key finding 2 — Solana validator targeting (ongoing)

SSH dictionaries continue to carry Solana-specific credentials: usernames `sol` (3,378), `solana` (2,837), `solv` (1,843); passwords `solana` (1,071), `sol` (1,034). A botnet is specifically hunting Solana validators/nodes exposing SSH — aiming at validator key theft or staking-reward redirection.

Key finding 3 — Multi-architecture Mirai/Gafgyt droppers

85 downloads, 33 unique SHA-256, classic multi-arch delivery. Four primary dropper servers:

Dropper	ASN / geo	Payload series
103[.]77[.]246[.]174	AS140810 Megacore, VN	full arch set + <code>/wget.sh</code>
45[.]81[.]234[.]64	AS44486 synlinq.de, DE	<code>/10Gbins.sh</code>
202[.]71[.]14[.]246	AS43641 Sollutium EU, NL	<code>proto.*</code>

Dropper	ASN / geo	Payload series
209[.]92[.]170[.]1225	AS212238 Datacamp, MY	lmkjn.*

Architecture-selection fingerprint command `/bin/./uname -s -v -n -r -m` (578 occurrences) precedes payload selection.

Key finding 4 — Captured malware: WannaCry/SMB dominance

126 unique samples, all malicious (sources: Dionaea 86, Cowrie 36, harvested-dropper 4). Approximately 85 of 126 are PE32 DLLs flagged by YARA `WannaCry_Ransomware + SUSP_Imphash_Mar23_2` — residual WannaCry/EternalBlue SMB traffic still scanning the internet. ELF samples (35) map to the Mirai/Gafgyt downloads above; SSH worms detected as `Go_SSHWorm_Generic` (5), `Panchan_SSHWorm_Go` (4), `Discord_C2_ELF` (5). Note: ClamAV produced **zero** detections across all 126 samples — current capture stream is detected only by YARA/imphash.

Key finding 5 — ICS substation decoy (IEC-104), preliminary

A recently deployed ICS decoy emulating a substation gateway/RTU (IEC-104) recorded, over a partial observation window, 205 IEC-104 events including 84 connects, **62 on-load tap-changer (OLTC) operations** and 28 general interrogations — i.e. attackers issued tap-changer commands against the emulated SCADA logic, not merely connected. Figures are preliminary given the short observation period.

Source operators and geography

5,542 operator hits (passive ASN-based attribution; not IOC-feed matches). Leading source operators: DigitalOcean (3,317), OVH (1,127), Alibaba Cloud (409), Amazon AWS (264), VPSVAULT.HOST (157), M247 (130). Leading source countries: US (11,132), CN (6,648), RO (3,714), NL (3,483), DE (2,845), RU (2,382). The bulk of attack volume originates from mainstream clouds (rented VPS botnets); sanctioned/bulletproof networks (Aeza, Stark brands) carry low volume but the targeted LLM-abuse load.

Botnet campaigns

SSH clustering observed campaigns of up to 670 unique IPs in a single cluster, covering up to ~1,900 IPs across as many as 53 parallel campaigns. Active C2 hosts include 14[.]46[.]136[.]77 (KR), 151[.]242[.]125[.]187 (HK), 87[.]121[.]79[.]73 (BG).

PL phishing sink

A phishing sink sensor recorded ~387,000 visits in May. Observed phishing-host clusters include brand impersonation of document-signing and security-awareness vendors and large sub-domain clusters; visit sources were predominantly cloud bots.

MITRE ATT&CK

T1595.002 (Active Scanning: Vulnerability Scanning) · T1110 (Brute Force) · T1046 (Network Service Discovery) · T1190 (Exploit Public-Facing Application) · T1610 (Deploy Container) · T1496 (Resource Hijacking — inference/compute theft) · T1552.005 (Cloud Instance Metadata API) · T1078.004 (Valid Accounts: Cloud) · T0855 (Unauthorized Command Message — ICS/IEC-104) · T0836 (Modify Parameter — OLTC/tap-changer operation, ICS).

IOCs

See `iocs.md` / machine-readable STIX 2.1 and MISP dumps. All indicators are defanged in human-readable documents and refanged in the JSON bundles.